

FEDERAL STATE AUTONOMOUS EDUCATIONAL INSTITUTION
OF HIGHER EDUCATION
ITMO UNIVERSITY

REPORT

on learning practice No.2

«Analysis of multivariate random variables»

Performed by:
Putnikov Semyon
Gabrielian Mikhail

St. Petersburg

2021

code

November 15, 2021

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sb
from scipy import stats as st
from sklearn.preprocessing import scale, PolynomialFeatures
from sklearn.decomposition import PCA
from sklearn.model_selection import train_test_split
from sklearn import linear_model as lm
from sklearn.metrics import mean_squared_error, r2_score
from tabulate import tabulate
import statsmodels.api as sm
```

```
[2]: source_data_path = "../tcs_stock.csv"
row_df = pd.read_csv(source_data_path)
row_df["Deliverble(%)"] = row_df['%Deliverble']
```

```
[3]: feature_cols = ['Open', 'Low', 'Last', 'Close', 'Trades', 'Deliverble(%)']
target_col = 'High'

df = row_df[feature_cols + [target_col]]
df.head()
```

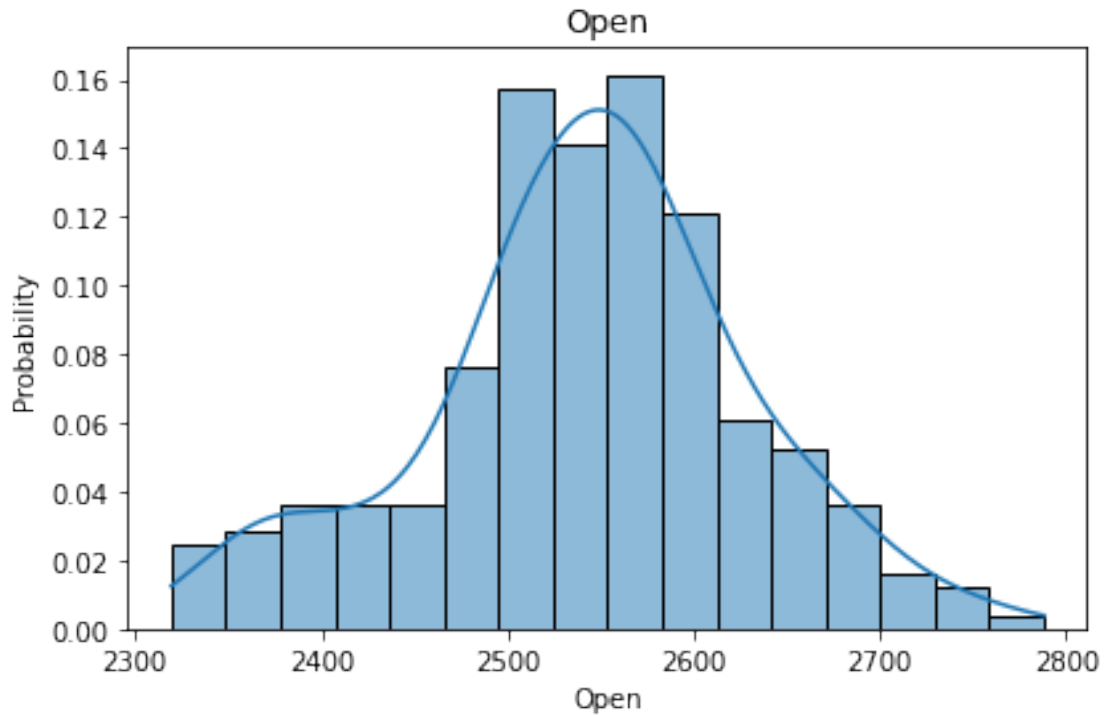
```
[3]:
```



	Open	Low	Last	Close	Trades	Deliverble(%)	High
0	2567.0	2541.00	2550.00	2545.55	8002	0.2883	2567.00
1	2551.0	2550.60	2588.40	2579.45	27585	0.6683	2590.95
2	2581.0	2524.65	2538.10	2540.25	43234	0.5207	2599.90
3	2529.1	2440.00	2450.05	2446.60	84503	0.5894	2529.10
4	2470.0	2407.45	2426.90	2417.70	101741	0.6724	2479.15

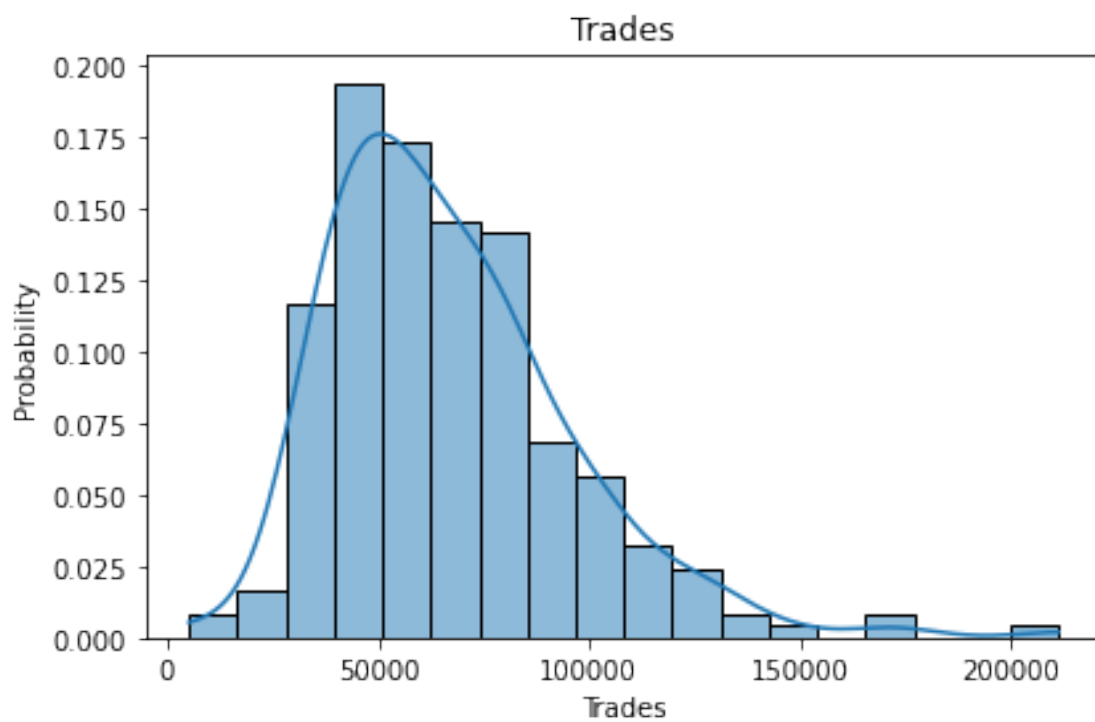

```
[4]: features = df[feature_cols]
target = df[target_col]
```


0.1 Step 1. You need to make a non-parametric estimation of PDF in form of histogram and using kernel density function for MRV (or probability law in case of discrete MRV)

```
[5]: for col in feature_cols:
      fig, ax = plt.subplots(tight_layout=True)
      sb.histplot(df[col], ax=ax, kde=True, stat="probability")
      plt.title(col)
      plt.show()
```









0.2 Step 2. You need to make an estimation of multivariate mathematical expectation and variance.

```
[6]: df.mean()
```

```
[6]: Open          2542.172782
Low             2514.408468
Last            2538.039718
Close           2537.717944
Trades          66873.608871
Deliverble(%)   0.670336
High            2563.580444
dtype: float64
```

```
[7]: df.var()
```


```
[7]: Open          7.674758e+03
Low             6.881163e+03
Last            7.542802e+03
Close           7.579063e+03
Trades          8.342223e+08
Deliverble(%)   8.275145e-03
High            8.208064e+03
```

dtype: float64

0.3 Step 3. You need to make a non-parametric estimation of conditional distributions, mathematical expectations and variances.

```
[8]: condition_more = df['Trades'] >= 66000  
condition_less = df['Trades'] < 66000
```

```
[9]: fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(20,8))  
  
sb.histplot(df[condition_more]['High'], ax=ax1, kde=True, stat="density")  
ax1.set_title("df['Trades'] >= 66000")  
  
sb.histplot(df[condition_less]['High'], ax=ax2, kde=True, stat="density")  
ax2.set_title("df['Trades'] < 66000")  
  
plt.show()
```



```
[10]: df[condition_more].mean()
```

```
[10]: Open          2558.012500  
Low            2523.619444  
Last           2550.018056  
Close          2549.429630  
Trades         92309.546296  
Deliverble(%)    0.691087  
High           2582.467130  
dtype: float64
```

```
[11]: df[condition_more].var()
```


```
[11]: Open          6.191144e+03
      Low           5.887515e+03
      Last          6.700395e+03
      Close         6.871748e+03
      Trades        5.982571e+08
      Deliverble(%) 6.254171e-03
      High          7.031163e+03
      dtype: float64
```

0.4 Step 4. You need to make an estimation of pair correlation coefficients, confidence intervals for them and significance levels.

```
[12]: fig, ax = plt.subplots(figsize=(8,8))

      corr = features.corr()
      mask = np.triu(np.ones_like(corr, dtype=bool))
      sb.heatmap(corr, mask=mask, annot=True, ax=ax, vmax=1, fmt='.5g', linewidths=.5)

      plt.title('Correlation')
      plt.show()
```

```
[13]: def _estimate_correlation(x, y):
        return st.pearsonr(x, y)

def _estimate_confidence_intervals(cor, x, y, alpha = 0.05):
    coeff = np.arctanh(cor)

    std = 1/np.sqrt(x.size-3)
    z = st.norm.ppf(1-alpha/2)
    return coeff-z*std, coeff+z*std
```

```
[14]: tab = [["Pair", "Correlation coeff", "Low border of conf interval", "High_
↪border of conf interval"]]
for i, x_col_name in enumerate(feature_cols):
    j = i + 1

    if j >= len(feature_cols):
        break

    for k in range(j, len(feature_cols)):
        y_col_name = feature_cols[k]
        x = df[x_col_name]
        y = df[y_col_name]

        cor, p = _estimate_correlation(x, y)

        low, high = _estimate_confidence_intervals(cor, x, y)
        tab.append([f'{x_col_name} - {y_col_name}', cor, low, high])
print(tabulate(tab, headers="firstrow", tablefmt="grid"))
```

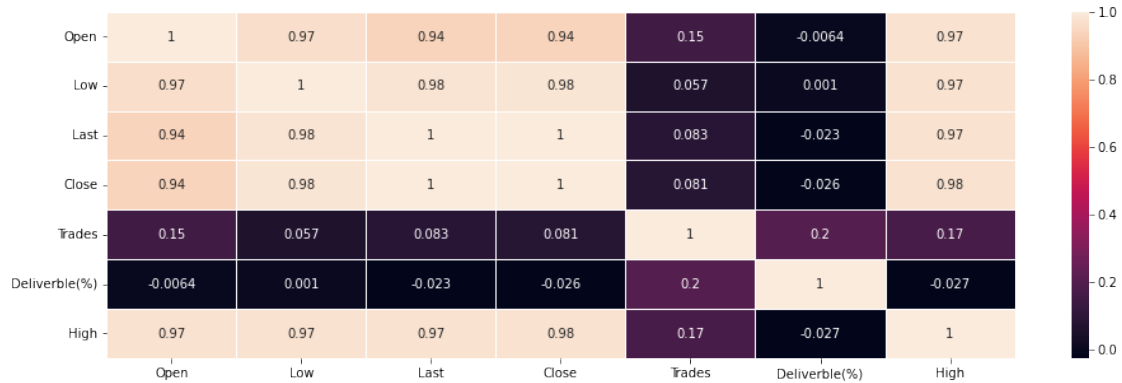
Pair	Correlation coeff	Low border of conf interval	High border of conf interval
Open - Low 2.17006	0.967063	1.91963	
Open - Last 1.82896	0.935876	1.57853	
Open - Close 1.84768	0.938159	1.59724	
Open - Trades 0.274698	0.148377	0.0242634	
Open - Deliverble(%) 0.118847	-0.00637008	-0.131588	
Low - Last 2.36535	0.977593	2.11492	

Low - Close	0.980559	2.18665
2.43708		
Low - Trades	0.0569236	-0.0682323
0.182203		
Low - Deliverble(%)	0.00101649	-0.124201
0.126234		
Last - Close	0.998138	3.36381
3.61424		
Last - Trades	0.0826734	-0.0423549
0.20808		
Last - Deliverble(%)	-0.0232686	-0.14849
0.101945		
Close - Trades	0.080853	-0.0441876
0.206247		
Close - Deliverble(%)	-0.0258261	-0.151049
0.0993856		
Trades - Deliverble(%)	0.201487	0.079064
0.329499		

0.5 Step 5. Choose a task formulation for regression. Estimate multivariate correlation (target -predictors).

Predict values of 'High' feature

```
[15]: plt.figure(figsize=(16, 5))
      sb.heatmap(df.corr(method='pearson'), annot=True, linewidths=.5)
      plt.show()
```



Use PCA method and see how many variables we need to take for regression.


```
[16]: #Standardize a dataset
std_df = scale(df)
std_df = pd.DataFrame(std_df, index=df.index, columns=df.columns)
```

```
[17]: std_df.head()
```

```
[17]:
```

	Open	Low	Last	Close	Trades	Deliverble(%)	High
0	0.283970	0.321211	0.137992	0.090146	-2.042407	-4.208179	0.037820
1	0.100965	0.437173	0.581031	0.480330	-1.363023	-0.022430	0.302709
2	0.444101	0.123712	0.000696	0.029144	-0.820119	-1.648263	0.401696
3	-0.149525	-0.898812	-1.015180	-1.048754	0.611609	-0.891524	-0.381355
4	-0.825503	-1.291997	-1.282272	-1.381388	1.209639	0.022732	-0.933805

```
[18]: pca = PCA().fit(std_df)
y = np.std(pca.transform(std_df), axis=0)**2
x = np.arange(len(y)) + 1
plt.plot(x, y, "o-")
plt.xticks(x, df.columns, rotation=60)
plt.ylabel("Variance")
plt.show()
```



I think that 3 features is ok. Let's use Open, Close, Low

0.6 Step 6. Build regression model and make an analysis of multicollinearity and regularization (if needed).

```
[33]: train, test = train_test_split(df[['High', 'Open', 'Close', 'Low']].copy())
```

```
train_feature = train[['Open', 'Close', 'Low']]
test_feature = test[['Open', 'Close', 'Low']]
```

```
train_target = train['High']
test_target = test['High']
```

```
[20]: tab = [["Type", "MSE", "R2", "Coeff"]]
```

```
regression = lm.LinearRegression()
regression.fit(train_feature, train_target)
predicted = regression.predict(test_feature)
tab.append(["Least Squares model", str(mean_squared_error(test_target,
    ↪ predicted)), str(r2_score(test_target, predicted)), str(regression.coef_)])
```

```
regression = lm.Lasso(alpha=0.0001, random_state=5)
```

```

regression.fit(train_feature, train_target)
predicted = regression.predict(test_feature)
tab.append(["Lasso model", str(mean_squared_error(test_target, predicted)),
↳str(r2_score(test_target, predicted)), str(regression.coef_)])

regression = lm.Ridge(random_state=1)
regression.fit(train_feature, train_target)
predicted = regression.predict(test_feature)
tab.append(["Ridge model", str(mean_squared_error(test_target, predicted)),
↳str(r2_score(test_target, predicted)), str(regression.coef_)])

print(tabulate(tab, headers="firstrow", tablefmt="grid"))

```

```

+-----+-----+-----+-----+
---+
| Type           |      MSE |      R2 | Coeff
|
+=====+=====+=====+=====+
===+
| Least Squares model | 154.324 | 0.973688 | [ 0.61183695  0.73816119
-0.32282928] |
+-----+-----+-----+-----+
---+
| Lasso model       | 154.324 | 0.973688 | [ 0.61183662  0.73816111
-0.32282885] |
+-----+-----+-----+-----+
---+
| Ridge model       | 154.325 | 0.973688 | [ 0.61182267  0.73813709 -0.3227891
] |
+-----+-----+-----+-----+
---+

```

0.7 Step 7. Analyze the quality of regression model (distribution of residuals, determination coefficient).

R^2 score can be find at the end of Step 6

```

[21]: X = df[['Open', 'Close', 'Low']]
      Y = df['High']

```

```

[22]: def _draw_qq_plot(predicted, predicted_all):
      percs = np.linspace(0, 100, 21)
      qn_first = np.percentile(predicted, percs)
      qn_second = np.percentile(predicted_all, percs)

      plt.figure(figsize=(8,8))

```

```

min_qn = np.min([qn_first.min(), qn_second.min()])
max_qn = np.min([qn_first.max(), qn_second.max()])
x = np.linspace(min_qn, max_qn)

plt.plot(qn_first, qn_second, ls="", marker="o", markersize=6)
plt.plot(x,x,color="k", ls="--")
plt.show

```

```


[34]: regression = lm.LinearRegression()
      regression.fit(train_feature, train_target)
      predicted = regression.predict(test_feature)

      print('R^2 score =', r2_score(test_target, predicted))

      predicted_all = regression.predict(X)
      _draw_qq_plot(predicted, predicted_all)


```

R^2 score = 0.9697279864152346



R^2 show a good quality of our model as a qqplot

```
[35]: residuals = Y - regression.predict(X)  
ax = sb.histplot(residuals, kde = 'True', stat="density")
```

```
[25]: st.kstest(residuals, 'norm', args=(residuals.mean(), residuals.var()))
```

```
[25]: KstestResult(statistic=0.4353543408175741, pvalue=2.886743745829084e-43)
```

Residuals are not distributed normally.

```
[26]: mod = sm.OLS(train_target, train_feature)
      res = mod.fit()
      print(res.conf_int(0.01))
```

	0	1
Open	0.510829	0.731068
Close	0.617991	0.902066
Low	-0.573751	-0.176412

```
[37]: residuals.describe()
```

```
[37]: count    248.000000
      mean      0.380230
      std     12.471677
      min    -24.312439
      25%     -8.113857
      50%     -0.971024
      75%      6.655066
      max     54.325712
```

Name: High, dtype: float64